

1.Introduction

[What's Machine Learning?](#)

[Supervised Learning](#)

[Unsupervised Learning](#)

2.Linear Regression with One Variable

[Model Representation](#)

[Cost Function](#)

[Gradient Descent](#)

3.Linear Algebra Review

[Matrices and Vectors](#)

[Matrix Multiplication Properties](#)

[Inverse and Transpose](#)

4.Linear Regression with Multiple Variables

[Multiple Features](#)

[Gradient Descent for Multiple Variables](#)

[Gradient Descent in Practice I - Feature Scaling](#)

[Gradient Descent in Practice II - Learning Rate](#)

[Features and Polynomial Regression](#)

[Normal Equation](#)

[Normal Equation and Non-invertibility](#)

5.Octave/Matlab Tutorial

[Basic Operations](#)

[Moving Data Around](#)

[Computing on Data](#)

[Plotting Data](#)

[Control Statements](#)

[Vectorization](#)

6.Logistic Regression

[Classification](#)

[Hypothesis Representation](#)

[Decision Boundary](#)

[Cost Function](#)

[Simplified Cost Function](#)

[Advanced Optimization](#)

[Multi-class Classification:One-vs-All](#)

7.Regularization

[The problem of overfitting](#)

[Cost Function](#)

[Regularized linear regression](#)

[Regularized logistic regression](#)

8.Neural Networks:Representation

[Non-linear hypotheses](#)

[Neurons and the brain](#)

[Model Representation](#)

[9.Neural Networks:Learning](#)

[Cost Function](#)

[Backpropagation algorithm](#)

[Backpropagation Intuition](#)

[Implementation Note:Unrolling Parameters](#)

[Gradient Checking](#)

[Random Initialization](#)

[Putting it together](#)

[10.Advice for Applying Machine Learning](#)

[Deciding What to Try Next](#)

[Evaluating a Hypothesis](#)

[Model Selection and Training/validation/test sets](#)

[Diagnosing Bias vs. Variance](#)

[Regularization and Bias/Variance](#)

[Learning Curves](#)

[Deciding What to Do Next](#)

[11.Machine Learning System Design](#)

[Prioritizing What to Work On](#)

[Error Analysis](#)

[Error Metrics for Skewed Classes](#)

[Trading Off Precision and Recall](#)

[Data For Machine Learning](#)

[12.Support Vector Machines](#)

[Optimization Objective](#)

[Large Margin Intuition](#)

[Mathematics Behind Large Margin Classification](#)

[Kernels](#)

[13.Unsupervised Learning](#)

[K-Means Algorithm](#)

[Optimization Objective](#)

[Random Initialization](#)

[Choosing the Number of Clusters](#)

[14.Dimensionality Reduction](#)

[Motivation I: Data Compression](#)

[Motivation II: Data Visualization](#)

[Principal Component Analysis Problem Formulation](#)

[Principal Component Analysis Algorithm](#)

[Choosing the Number of Principal Components](#)

[Reconstruction from Compressed Representation](#)

[Advice for Applying PCA](#)

[15.Anomaly Detection](#)

[Problem Motivation](#)

[Gaussian Distribution](#)

[Algorithm](#)

[Developing and Evaluating an anomaly detection system](#)

[Anomaly Detection vs. Supervised Learning](#)

[Choosing What Features to Use](#)

[Multivariate Gaussian Distribution](#)

[Anomaly Detection using the Multivariate Gaussian Distribution](#)

16.Recommender Systems

[Problem formulation](#)

[Content Based Recommendations](#)

[Collaborative filtering](#)

[Collaborative filtering algorithm](#)

[Vectorization: Low Rank Matrix Factorization](#)

[Implementation Detail: Mean Normalization](#)

17.Large Scale Machine Learning

[Learning with Large Datasets](#)

[Stochastic Gradient Descent](#)

[Mini-Batch Gradient Descent](#)

[Stochastic Gradient Descent Convergence](#)

[Online Learning](#)

[Map-reduce and data parallelism](#)

18.Application Example: Photo OCR

[Problem Description and Pipeline](#)

[Sliding Windows](#)

[Getting Lots of Data and Artificial Data Synthesis](#)

[Caption Analysis: What Part of the Pipeline to Work on Next](#)

1.Introduction

What's Machine Learning?

- Arthur Samuel's definition

Field of study that gives computers the ability to learn without being explicitly programmed.

- Tom Mitchell's definition

A computer program is said to learn from experience E with respect to some task T and some performance measure P , if its performance on T , as measured by P , improves with experience E .

- Classifications:
 - Supervised Learning
 - Unsupervised Learning

Supervised Learning

- definition

we are given a data set and already know what our correct output should look like, having the idea that there is a relationship between the input and the output

- problems

- regression

we are trying to predict results within a continuous output, meaning that we are trying to map input variables to some continuous function.

- classification

we are instead trying to predict results in a discrete output.

we are trying to map input variables into discrete categories.

Unsupervised Learning

- definition

unsupervised learning allows us to approach problems with little or no idea what our results should look like. We can derive structure from data where we don't necessarily know of the variables.

We can derive this structure by clustering the data based on relationships among the variables in the data.

2.Linear Regression with One Variable

Model Representation

- notation

$(x^{(i)}, y^{(i)})$: training example . input-output or input_features-target

m : a training set

X / Y : the space of input/output values . $X = Y = \mathbb{R}$

$h(x)$: a predictor for the corresponding value of y . called a hypothesis

eg. $h_{\theta}(x) = \theta_0 + \theta_1 x$

- problem
 - regression problem
when the target variable that we're trying to predict is continuous.
 - classification problem
when y can take on only a small number of discrete values.

Cost Function

- function
 - hypothesis function
 $h_{\theta}(x) = \theta_0 + \theta_1 x$
 - parameters
 θ_0, θ_1
 - cost function / mean squared error function
 $J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i)^2$
 - goal
 $\underset{\theta_0, \theta_1}{\text{minimize}} J(\theta_0, \theta_1)$

Gradient Descent

- Outline:
 - start with some parameters
common choice : $\theta_0 = 0, \theta_1 = 0$
 - keep changing θ_0, θ_1 to reduce $J(\theta_0, \theta_1)$
until we hopefully end up at a local minimum

- algorithm

repeat simultaneously until convergence:

```
{  $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$  (for  $j = 0$  and  $j = 1$ )  
}
```

α : learning rate. controls how big a step

- α is too small: gradient descent can be slow
- α is too large: gradient descent can overshoot the minimum.

It may fail to converge, or even diverge.

- batch

batch gradient descent: refers to the fact that in every step of gradient descent, we're looking at all of the training examples.

3. Linear Algebra Review

Matrices and Vectors

Matrix : Rectangular array of numbers

Vector : An $n \times 1$ matrix

Matrix Multiplication Properties

$$A * B \neq B * A$$

$$(A * B) * C = A * (B * C)$$

Identity Matrix : denoted $I_{n \times n}$

$$AI = IA = A$$

Inverse and Transpose

$$\text{inverse of } A = \text{inv}(A)$$

$$\text{transpose of } A = A'$$

4. Linear Regression with Multiple Variables

Multiple Features

n : number of features

$x^{(i)}$: input features of i^{th} training example (\mathbb{R}^n)

$x_j^{(i)}$: value of feature j in i^{th} training example

Gradient Descent for Multiple Variables

Hypothesis : $h_{\theta}(x) = \theta^T x$

Parameters : $\theta = [\theta_0 \ \theta_1 \dots \theta_n]$

Cost function: $J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$

Gradient descent:

$$\text{Repeat} \{ \theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \}$$

$$\text{Repeat} \{ \theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) * x_j^{(i)} \}$$

Gradient Descent in Practice I - Feature Scaling

Idea: Make sure features are on a similar scale.

Get every feature into approximately a $-1 \leq x_i \leq 1$ range.

- Mean normalization

replace x_i with $x_i - \mu_i$ to make features have approximately 0 mean

$$x_1 := \frac{x_1 - \mu_1}{s_1}$$

μ_1 : average value of x_1 in training set

s_1 : range(max - min)

Gradient Descent in Practice II - Learning Rate

if α is too small : slow convergence.

if α is too large : $J(\theta)$ may not decrease on every iteration; may not converge.

Features and Polynomial Regression

We can change the behavior or curve of our hypothesis function by making it a quadratic, cubic or square root function (or any other form).

Normal Equation

To compute θ , and then minimize J function:

$$\theta = (X^T X)^{-1} X^T y$$

Gradient Descent	Normal Equation
1. Need to choose α	1. No need to choose α
2. Need many iterations	2. Do not need to iterate
3. $O(kn^2)$ Work well even when n is large	3. Need to compute $(X^T X)^{-1}$
	4. $O(n^3)$ Slow if n is very large

Normal Equation and Non-invertibility

`inv()`: $(X^T X)$ can't be non-invertible

`pinv()`: $(X^T X)$ even be non-invertible (pseudo-inverse)

Noninvertible's causes:

1. Redundant features, where two features are very closely related (i.e. linearly dependent)
2. Too many features (i.e. $m \leq n$). We should delete some features or use regularization.

5. Octave/Matlab Tutorial

Basic Operations

```
>> a = pi
a = 3.1416
>> disp(sprintf('show 3 decimals: %0.3f', a))
show 3 decimals: 3.142
>> format long % change the format of value
>> a
a = 3.14159265358979
```



```
>> a = 1:0.3:2
a =
    1.0000    1.3000    1.6000    1.9000
```

```
>> a = zeros(1,3)
a =
    0    0    0
```

```
>> a = 3*ones(2,4)
a =
    3    3    3    3
    3    3    3    3
```

```
>> a=rand(2,3)
a =
    0.25599    0.54999    0.34204
    0.73561    0.31876    0.76891
```

```
>> a = randn(2,3) % Gaussian distribution with mean=0
a =
    0.32532   -0.26668    0.15293
    0.29101   -0.25914   -0.36278
```

```
>> a = 5 + 4*randn(1000,1);
>> hist(a) % get the pic
```

```
>> a = eye(3)
a =
Diagonal Matrix
    1    0    0
    0    1    0
    0    0    1
```

Moving Data Around

```
>> a = [1 2 ;3 4;5 6]
```

```
a =
```

```
    1    2
```

```
    3    4
```

```
    5    6
```

```
>> a(2,2)
```

```
ans = 4
```

```
>> a(3,:)
```

```
ans =
```

```
    5    6
```

```
>> a=[a,[11;12;13]]
```

```
a =
```

```
    1    2   11
```

```
    3    4   12
```

```
    5    6   13
```

```
>> size(a,1)
```

```
ans = 3
```

```
>> size(a,2)
```

```
ans = 2
```

```
>> length(a)
```

```
ans = 3
```

```
>> pwd
```

```
ans = /home/yule
```

```
>> who
```

```
Variables in the current scope:
```

```
a    ans  b    sz
```

```
>> whos
```

```
Variables in the current scope:
```

Attr	Name	Size	Bytes	Class
====	====	====	=====	=====
	a	3x2	48	double
	ans	1x10	10	char
	b	1x2	2	char
	sz	1x2	16	double

```
Total is 20 elements using 76 bytes
```

```
>> clear
```

```
>> load featuresX.dat
>> load('featuresX.dat')
>> featuresX %show the data
>> size(featuresX)
ans =
    47     2
>> v = featuresX(1:10)
>> save save_v.mat v
>> save save_v.txt v
>> load save_v.mat
```

```
>> a=[1 2;3 4]
a =
     1     2
     3     4
>> a(:)
ans =
     1
     3
     2
     4
```

Computing on Data

```
>> a
a =

     1     2
     3     4
     5     6
>> a >3
ans =

     0     0
     0     1
     1     1
>> find(a>3)
ans =

     3
     5
     6
>> max(a) % <=> max(a,[],1)
ans =

     5     6
>> max(a,[],2)
ans =

     2
     4
     6
>> max(max(a))
ans = 6
```

```

>> a=magic(3)
a =
     8     1     6
     3     5     7
     4     9     2
>> sum(a,1)
ans =
    15    15    15
>> sum(a,2)
ans =
    15
    15
    15
>> b=a.*eye(3)
b =
     8     0     0
     0     5     0
     0     0     2
>> b'
ans =
     8     0     0
     0     5     0
     0     0     2
>> flipud(b)
ans =
     0     0     2
     0     5     0
     8     0     0

```

Plotting Data

```

>> help plot
>> subplot(1,2,1)
>> axis([0 1 -1 1])
>> imagesc(a)
>> imagesc(a),colormap gray;

```

Control Statements

```

>> for i=1:10,
v(i)=2^i;
end;

```

```
function [y1,y2] = returntwo(x)
y1=x^2;
y2=x^3;
>> [a,b] = returntwo(9)
a = 81
b = 729
```

Vectorization

$$\sum_{j=0}^n \theta_j x_j = \theta^T x$$

6.Logistic Regression

Classification

Logistic Regression: $0 \leq h_{\theta}(x) \leq 1$

$h_{\theta}(x)$: estimated probability that $y=1$ on input x

Hypothesis Representation

$$h_{\theta}(x) = g(\theta^T x)$$

$$g(z) = \frac{1}{1 + e^{-z}}$$

Decision Boundary

Suppose predict ' $y=0$ ' if $h_{\theta}(x) \leq 0.5 : \theta^T x \leq 0$

Cost Function

cost function : $Cost(h_{\theta}(x), y) = \frac{1}{2}(h_{\theta}(x) - y)^2$ (non-convex)

$$Cost(h_{\theta}(x), y) = -\log(h_{\theta}(x)) \text{ (if } y = 1)$$

$$Cost(h_{\theta}(x), y) = -\log(1 - h_{\theta}(x)) \text{ (if } y = 0)$$

$$Cost(h_{\theta}(x), y) = -y \log(h_{\theta}(x)) - (1 - y) \log(1 - h_{\theta}(x))$$

Simplified Cost Function

$$J(\theta) = -\frac{1}{m} [\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))]$$

To fit parameters θ :

Want $\min_{\theta} J(\theta)$:

Repeat : $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$

Advanced Optimization

Introduction of other optimization method.

Multi-class Classification:One-vs-All

Easy to understand.

7.Regularization

The problem of overfitting

overfitting: high variance

Options:

1. Reduce number of features
 1. Manually select which features to keep
 2. Model selection algorithm
2. Regularization
 1. Keep all the features, but reduce magnitude/values of parameters θ_j
 2. Works well when we have a lot of features, each of which contributes a bit to predicting y .

Cost Function

Aim - Small values for parameters :

1. simpler hypothesis
2. less prone to overfitting

$$J(\theta) = \frac{1}{m} \left[\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

then $\min_{\theta} J(\theta)$

Regularized linear regression

1. Gradient descent

Repeat {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$\theta_j := \theta_j - \alpha \left[\left(\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \right) + \frac{\lambda}{m} \theta_j \right], \quad j \in \{1, 2, \dots, n\}$$

}

也可以移项得到更新表达式的另一种表示形式

$$\theta_j := \theta_j (1 - \alpha \frac{\lambda}{m}) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

2. Normal equation

$$\theta = (X^T X + \lambda \cdot L)^{-1} X^T y$$

$$\text{where } L = \begin{bmatrix} 0 & & & & \\ & 1 & & & \\ & & 1 & & \\ & & & \ddots & \\ & & & & 1 \end{bmatrix}$$

Regularized logistic regression

Cost function:

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

Gradient descent:

Repeat {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$\theta_j := \theta_j - \alpha \left[\left(\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \right) + \frac{\lambda}{m} \theta_j \right], \quad j \in \{1, 2, \dots, n\}$$

}

8. Neural Networks: Representation

Non-linear hypotheses

Neurons and the brain

Model Representation

Nothing worths recording.

9. Neural Networks: Learning

Cost Function

- preview: Cost function of Logistic Regression:

$$h_{\theta}(x) = g(\theta^T x) = z$$

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

- **Cost Function of Neural Network:**

$$h_{\Theta}(x) \in \mathbb{R}^K \quad (h_{\Theta}(x))_i = i^{th} \text{ output}$$

L = total no. of layers in network

s_l = no. of units in layer

$$J(\Theta) = -\frac{1}{m} \left[\sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(h_{\Theta}(x^{(i)}))_k + (1 - y_k^{(i)}) \log(1 - (h_{\Theta}(x^{(i)}))_k) \right] + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (\Theta_{ji}^{(l)})^2$$

Backpropagation algorithm

- **Forward propagation** : a_j^l : the activation of node j in layer l

$$a^{(1)} = x$$

$$z^{(i)} = \Theta^{(i-1)} a^{(i-1)}$$

$$a^{(i)} = g(z^{(i)}) \text{ (add } a_0^{(i)})$$

- **Back Propagation** : δ_j^l : the 'error' of node j in layer l

$$\delta^{(L)} = a^{(L)} - y$$

$$\delta^{(i)} = (\Theta^{(i)})^T \delta^{(i+1)} \cdot g'(z^{(i)})$$

$$g'(z^{(i)}) = a^{(i)} \cdot (1 - a^{(i)})$$

Formally,

$$\text{cost}(i) = y^{(i)} \log h_{\Theta}(x^{(i)}) + (1 - y^{(i)}) \log h_{\Theta}(x^{(i)})$$

$$\delta_j^{(l)} = \frac{\partial}{\partial z_j^{(l)}} \text{cost}(i)$$

- **Backpropagation algorithm**

Aim: to $\min_{\Theta} J(\Theta)$

$$\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta) = a_j^{(l)} \delta_i^{(l+1)} \quad (\lambda = 0)$$

Training set $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}) \dots (x^{(m)}, y^{(m)})\}$

Set $\Delta_{ij}^{(l)} = 0$

For $i=1$ to m :

Forward_propagation(x, Θ) #compute $a^{(l)}$

Back_Propagation(a, y, Θ) #compute $\delta^{(l)}$

$$\Delta_{ij}^{(l)} := \Delta_{ij}^{(l)} + a_j^{(l)} \delta_i^{(l+1)}$$

if $j == 0$:

$$\Delta_{ij}^{(l)} := \frac{1}{m} \Delta_{ij}^{(l)}$$

if $j \neq 0$:

$$\Delta_{ij}^{(l)} := \frac{1}{m} \Delta_{ij}^{(l)} + \lambda \Theta_{ij}^{(l)}$$

$$\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta) = \Delta_{ij}^{(l)}$$

Backpropagation Intuition

$\delta(i) = (\Theta^{(i)})^T \delta^{i+1}$ #I don't know which is correct...

Implementation Note: Unrolling Parameters

Gradient Checking

Random Initialization

Symmetry breaking:

Initialize each $\Theta_{ij}^{(l)}$ to a random value in $[-\epsilon, \epsilon]$

(i.e. $-\epsilon \leq \Theta_{ij}^{(l)} \leq \epsilon$)

Putting it together

- Pick a network architecture
 1. reasonable default: 1 hidden layer
 2. if hidden layer > 1, have same no. of hidden units in every layer
 3. usually the more the better
- Training a neural network
 1. Randomly initialize weights
 2. Implement forward propagation to get $h_{\Theta}(x^{(i)})$ for any $x^{(i)}$
 3. Implement code to compute cost function $J(\Theta)$
 4. Implement backprop to compute partial derivatives $\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta)$
 5. Use gradient checking to compare $\frac{\partial}{\partial \Theta_{jk}^{(l)}} J(\Theta)$ computed using backpropagation vs. using numerical estimate of gradient of $J(\Theta)$. Then disable gradient checking code.
 6. Use gradient descent or advanced optimization method with backpropagation to try to minimize $J(\Theta)$ as a function of parameters Θ

10. Advice for Applying Machine Learning

Deciding What to Try Next

- Next:
 1. Get more training examples

2. Try smaller sets of features
3. Try getting additional features
4. Try adding polynomial features
5. Try decreasing λ
6. Try increasing λ

- Machine learning diagnostic

A test that you can run to gain insight what is or not working with a learning algorithm, and gain guidance as to how best to improve its performance.

Evaluating a Hypothesis

```
if((h(x)>=0.5)&&(y==0) || (h(x)<=0.5)&&(y==1))
    err(h(x),y) = 1
else
    err(h(x),y) = 0
```

$$\text{Test error} = \frac{1}{m_{\text{test}}} \sum_{i=1}^{m_{\text{test}}} \text{err}(h(x_{\text{test}}^{(i)}), y^{(i)})$$

Model Selection and Training/validation/test sets

60% : 20% : 20%

Diagnosing Bias vs. Variance

High bias is underfitting and high variance is overfitting. Ideally, we need to find a golden mean between these two.

High bias (underfitting): both $J_{\text{train}}(\Theta)$ and $J_{\text{CV}}(\Theta)$ will be high. Also, $J_{\text{CV}}(\Theta) \approx J_{\text{train}}(\Theta)$.

High variance (overfitting): $J_{\text{train}}(\Theta)$ will be low and $J_{\text{CV}}(\Theta)$ will be much greater than $J_{\text{train}}(\Theta)$.

Regularization and Bias/Variance

choose the model and the regularization term λ :

1. Create a list of lambdas (i.e. $\lambda \in \{0, 0.01, 0.02, 0.04, 0.08, 0.16, 0.32, 0.64, 1.28, 2.56, 5.12, 10.24\}$);
2. Create a set of models with different degrees or any other variants.

3. Iterate through the λ s and for each λ go through all the models to learn some Θ .
4. Compute the cross validation error using the learned Θ (computed with λ) on the $J_{CV}(\Theta)$ **without** regularization or $\lambda = 0$.
5. Select the best combo that produces the lowest error on the cross validation set.
6. Using the best combo Θ and λ , apply it on $J_{test}(\Theta)$ to see if it has a good generalization of the problem.

Learning Curves

- **Experiencing high bias:**

Low training set size: causes $J_{train}(\Theta)$ to be low and $J_{CV}(\Theta)$ to be high.

Large training set size: causes both $J_{train}(\Theta)$ and $J_{CV}(\Theta)$ to be high with $J_{train}(\Theta) \approx J_{CV}(\Theta)$.

If a learning algorithm is suffering from **high bias**, getting more training data will not **(by itself)** help much.

- **Experiencing high variance:**

Low training set size: $J_{train}(\Theta)$ will be low and $J_{CV}(\Theta)$ will be high.

Large training set size: $J_{train}(\Theta)$ increases with training set size and $J_{CV}(\Theta)$ continues to decrease without leveling off. Also, $J_{train}(\Theta) < J_{CV}(\Theta)$ but the difference between them remains significant.

If a learning algorithm is suffering from **high variance**, getting more training data is likely to help.

Deciding What to Do Next

Our decision process can be broken down as follows:

- **Getting more training examples:** Fixes high variance
- **Trying smaller sets of features:** Fixes high variance
- **Adding features:** Fixes high bias
- **Adding polynomial features:** Fixes high bias
- **Decreasing λ :** Fixes high bias
- **Increasing λ :** Fixes high variance.

Model Complexity Effects:

- Lower-order polynomials (low model complexity) have high bias and low variance. In this case, the model fits poorly consistently.
- Higher-order polynomials (high model complexity) fit the training data extremely well and the test data extremely poorly. These have low bias on the training data, but very high variance.
- In reality, we would want to choose a model somewhere in between, that can generalize well but also fits the data reasonably well.

11. Machine Learning System Design

Prioritizing What to Work On

how to improve the accuracy of this classifier:

- Collect lots of data (for example "honeypot" project but doesn't always work)
- Develop sophisticated features (for example: using email header data in spam emails)
- Develop algorithms to process your input in different ways (recognizing misspellings in spam).

Error Analysis

The recommended approach to solving machine learning problems is to:

- Start with a simple algorithm, implement it quickly, and test it early on your cross validation data.
- Plot learning curves to decide if more data, more features, etc. are likely to help.
- Manually examine the errors on examples in the cross validation set and try to spot a trend where most of the errors were made.

Error Metrics for Skewed Classes

- Precision/Recall

$$\begin{aligned} \circ \text{Precision} &= \frac{\text{True positive}}{\text{True positive} + \text{False positive}} \\ \circ \text{Recall} &= \frac{\text{True positive}}{\text{True positive} + \text{False negative}} \end{aligned}$$

Trading Off Precision and Recall

1. if we want to predict $y=1$ only if very confident:
Higher precision, lower recall
2. if we want to avoid missing too many cases of $y=1$

Lower precision, Higher recall

$$F_1 \text{ Score} = 2 \frac{\text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}}$$

Data For Machine Learning

12.Support Vector Machines

Optimization Objective

- Logistics Regression

$$\min_{\theta} \frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \text{cost}_1(\theta^T x^{(i)}) + (1 - y^{(i)}) \text{cost}_0(\theta^T x^{(i)}) \right] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

$$\text{cost}_1(\theta^T x^{(i)}) = -\log h_{\theta}(x^{(i)})$$

$$\text{cost}_0(\theta^T x^{(i)}) = -\log(1 - h_{\theta}(x^{(i)}))$$

- Support Vector Machine

1. Cost Function

$$\min_{\theta} C \left[\sum_{i=1}^m y^{(i)} \text{cost}_1(\theta^T x^{(i)}) + (1 - y^{(i)}) \text{cost}_0(\theta^T x^{(i)}) \right] + \frac{1}{2} \sum_{j=1}^n \theta_j^2$$

$$\text{cost}_1(\theta^T x^{(i)}) = -\log h_{\theta}(x^{(i)})$$

$$\text{cost}_0(\theta^T x^{(i)}) = -\log(1 - h_{\theta}(x^{(i)}))$$

$$C = \frac{1}{\lambda}$$

Large C: Lower bias, high variance

Small C: Higher bias, low variance

2. Hypothesis

$$\text{if } \theta^T x \geq 0 : h_{\theta}(x) = 1, \text{ otherwise : } h_{\theta}(x) = 0$$

Large Margin Intuition

if $y = 1$, we want $\theta^T x \geq 1$

if $y = 0$, we want $\theta^T x \leq -1$

Mathematics Behind Large Margin Classification

- Vector Inner Product

$$u = \begin{bmatrix} u_1 \\ u_2 \end{bmatrix} v = \begin{bmatrix} v_1 \\ v_2 \end{bmatrix}$$

$$\|u\| = \sqrt{u_1^2 + u_2^2}$$

$$\text{projection: } p = \frac{u^T v}{\|u\|}$$

- SVM Decision Boundary

When C is very large, the cost function is $\min_{\theta} [C * 0 + \frac{1}{2} \sum_{j=1}^n \theta_j^2]$

simplification: C is very large, $\theta_0 = 0$, $n = 2$

$$\min_{\theta} : \frac{1}{2} \sum_{j=1}^n \theta_j^2 = \frac{1}{2} (\theta_1^2 + \theta_2^2) = \frac{1}{2} (\sqrt{\theta_1^2 + \theta_2^2})^2 = \frac{1}{2} \|\theta\|^2$$

$$\text{s.t. } \theta^T x^{(i)} \geq 1 \quad \text{if } y^{(i)} = 1$$

$$\theta^T x^{(i)} \leq -1 \quad \text{if } y^{(i)} = 0$$

$$\iff p^{(i)} \geq 1 \quad \text{if } y^{(i)} = 1$$

$$p^{(i)} \leq -1 \quad \text{if } y^{(i)} = 0$$

where $p^{(i)}$ is the projection of $x^{(i)}$ onto the vector θ

Kernels

- Gaussian kernel function

$$k(x, l^{(i)}) = \text{similarity}(x, l^{(i)}) = \exp\left(-\frac{\|x - l^{(i)}\|^2}{2\sigma^2}\right)$$

1. if $x \approx l^{(i)}$:

$$k(x, l^{(i)}) \approx 1$$

2. if x is far from l:

$$k(x, l^{(i)}) \approx 0$$

Large σ^2 : Features f_i vary more smoothly. Higher bias, lower variance.

Small σ^2 : Features f_i vary less smoothly. Lower bias, higher variance.

- SVM with Kernels

Given $(x^{(2)}, y^2), (x^{(2)}, y^2) \dots, (x^{(m)}, y^m)$

choose landmarks : $l^{(1)} = x^{(1)}, l^{(2)} = x^{(2)} \dots, l^{(m)} = x^{(m)}$

Given example x, compute features f:

$$f = \begin{bmatrix} f_0 = 1 \\ f_1 = \text{similarity}(x, l^{(1)}) \\ f_2 = \text{similarity}(x, l^{(2)}) \\ \dots \\ f_m = \text{similarity}(x, l^{(m)}) \end{bmatrix}$$

- Hypothesis:

Predict "y=1" if $\theta^T f \geq 0$

- Training:

$$\min_{\theta} C[\sum_{i=1}^m y^{(i)} \text{cost}_1(\theta^T f^{(i)}) + (1 - y^{(i)}) \text{cost}_0(\theta^T f^{(i)})] + \frac{1}{2} \sum_{j=1}^m \theta_j^2$$

- Using an SVM

package: `liblinear` `libsvm`

Step:

1. choose parameter C
2. choose kernel
3. choose σ^2

13. Unsupervised Learning

K-Means Algorithm

Step1. cluster assignment

Step2. move centroid

```
Randomly initialize K cluster centroids u1,u2,...uK
Repeat{
  #step 1
  for i = 1 to m:
    c := index(from 1 to K) of cluster centroid closest to x
    # c = min(x-uk)^2
  #step 2
  for k = 1 to K:
    uk := average(mean) of points assigned to cluster k
}
```

Optimization Objective

$c^{(i)}$: index of cluster(1,2,...,K) to which example $x^{(i)}$ is currently assigned

μ_k : cluster centroid k ($\mu_k \in \mathbb{R}^n$)

$\mu_{c^{(i)}}$: cluster centroid of cluster to which example $x^{(i)}$ has been assigned

Optimization objective:

distortion cost function(失真代价函数):

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$
$$\min_{c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

Random Initialization

Randomly pick K training examples: Set μ_1, \dots, μ_K equal to these K examples

```
for i = 1 to 100 { # to avoid local optima
    Randomly initialize K-means
    Run K-means.
    Compute distortion cost function J()
}
```

Choosing the Number of Clusters

- Choosing the value of K:

Anyway...

14. Dimensionality Reduction

Motivation I: Data Compression

Motivation II: Data Visualization

Principal Component Analysis Problem Formulation

Reduce from n-dimension to k-dimension:

Find k vectors onto which to project the data, so as to minimize the projection error.

PS. PCA is not linear regression

Principal Component Analysis Algorithm

- Data preprocessing

1. mean normalization:

$$\mu_j = \frac{1}{m} \sum_{i=1}^m x_j^{(i)}$$

Replace each $x_j^{(i)}$ with $x_j - \mu_j$

2. scale features:

$$x_j = \frac{x_j - \mu_j}{s_j}$$

- PCA algorithm

1. Compute covariance matrix:

$$\Sigma = \frac{1}{m} \sum_{i=1}^n (x^{(i)})(x^{(i)})^T$$

2. Compute eigenvectors of matrix Σ :

$[U, S, V] = \text{svd}(\Sigma)$ or use $\text{eig}(\Sigma)$ # Singular value decomposition

where $U = [u^{(1)}, u^{(2)}, \dots, u^{(n)}] \in R^{n \times n}$

so $U_{\text{reduce}} = [u^{(1)}, u^{(2)}, \dots, u^{(k)}] \in R^{n \times k}$

$$Z = U_{\text{reduce}}^T X \in R^k$$

Choosing the Number of Principa Components

Average squared projection error:

$$\frac{1}{m} \sum_{i=1}^m \|x^{(i)} - x_{\text{approx}}^{(i)}\|^2$$

Total variation in the data:

$$\frac{1}{m} \sum_{i=1}^m \|x^{(i)} - x_{\text{approx}}^{(i)}\|^2$$

choose k to be smallest value so that

$$\frac{\frac{1}{m} \sum_{i=1}^m \|x^{(i)} - x_{\text{approx}}^{(i)}\|^2}{\frac{1}{m} \sum_{i=1}^m \|x^{(i)} - x_{\text{approx}}^{(i)}\|^2} \leq 1\%$$

(It means 99% of variance is retained)

For given k:

$$[U, S, V] = \text{svd}(\Sigma) \text{ where } S = \begin{bmatrix} S_{11} & & & \\ & S_{22} & & \\ & & \dots & \\ & & & S_{nn} \end{bmatrix}$$

$$\frac{\frac{1}{m} \sum_{i=1}^m \|x^{(i)} - x_{\text{approx}}\|^2}{\frac{1}{m} \sum_{i=1}^m \|x^{(i)} - x_{\text{approx}}\|^2} = 1 - \frac{\sum_{i=1}^k S_{ii}}{\sum_{i=1}^n S_{ii}}$$

So pick smallest value of k for which $\frac{\sum_{i=1}^k S_{ii}}{\sum_{i=1}^n S_{ii}} \geq 99\%$

Reconstruction from Compressed Representation

$$Z \in R^k \quad U_{\text{reduce}} \in R^{n \times k}$$

$$X_{\text{approx}} = U_{\text{reduce}} Z \in R^n$$

Advice for Applying PCA

Before implementing PCA, first try running whatever you want to do with the original data $x^{(i)}$.

Only if that doesn't do what you want, then implement PCA and consider using $z^{(i)}$.

15. Anomaly Detection

Problem Motivation

Example:

1. Fraud detection
2. Manufacturing
3. Monitoring computers in a data center

Gaussian Distribution

$$p(x; \mu; \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

- Parameter estimation

$$\text{mean: } \mu = \frac{1}{m} \sum_{i=1}^m x^{(i)}$$

variance: $\sigma^2 = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu)^2$

standard deviation: σ

Algorithm

Given new example x , compute $p(x)$:

$$p(x) = \prod_{j=1}^n p(x_j; \mu_j, \sigma_j^2)$$

Anomaly if $p(x) < \epsilon$

Developing and Evaluating an anomaly detection system

1. Fit model $p(x)$ on training set
2. On a cross validation/test example x , predict

$$f(x) = \begin{cases} 1 & \text{if } p(x) < \epsilon \text{ (anomaly)} \\ 0 & \text{if } p(x) \geq \epsilon \text{ (normal)} \end{cases}$$

3. Possible evaluation metrics:
 - True positive, false positive, false negative, true negative
 - Precision/Recall
 - F1-score
4. use cross validation set to choose parameter ϵ

Anomaly Detection vs. Supervised Learning

Anomaly Detecton	Supervised Learning
Very small number of positive examples.	Large number of positive and negative examples.
Hard for any algorithm to learn from positive examples what the anomalies look like.	Enough positive examples for algorithm to get a sense of what positive examples are like.
future anomalies may look nothing like any of the anomalous examples we've seen so far.	Future positive examples likely to be similar to ones in training set.

Choosing What Features to Use

Choose features that might take on unusually large or small values in the event of an anomaly.

Multivariate Gaussian Distribution

Don't model $p(x_1), p(x_2), \dots$, *tec.* separately.

Model $p(x)$ all in one go.

Parameters: $\mu \in \mathbb{R}^n, \Sigma \in \mathbb{R}^{n \times n}$

$$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp\left[-\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu)\right]$$

Anomaly Detection using the Multivariate Gaussian Distribution

1. Fit model $p(x)$ by setting:

$$\text{mean: } \mu = \frac{1}{m} \sum_{i=1}^m x^{(i)}$$

$$\text{variance: } \Sigma = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu)(x^{(i)} - \mu)^T$$

2. Given a new example x , compute:

$$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp\left[-\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu)\right]$$

Flag an anomaly if $p(x) < \epsilon$

Original model	Multivariate Gaussian
$p(x_1; \mu_1, \sigma_1^2) \times \dots \times p(x_n; \mu_n, \sigma_n^2)$	$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{n/2} \Sigma ^{1/2}} \exp\left[-\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu)\right]$
Manually create features to capture anomalies where x_1, x_2 take unusual combinations of values.	Automatically captures correlations between features.
Computationally cheaper	computationally more expensive
OK even if m is small. (m : training set size)	Must have $m > n$, or else Σ is non-invertible

16. Recommender Systems

Problem formulaiton

n_u : no.users

n_m : no.movies

$r(i, j)$: 1 if user j has rated movie i

$y^{(i,j)}$: rating given by user j to movie i

Content Based Recommendations

$\theta^{(j)}$: parameter vector for user j

$x^{(i)}$: feature vector for movie i

$m^{(j)}$: no. of movies rated by user j

For user j, movie i, predicted rating : $(\theta^{(j)})^T (x^{(i)})$

To learn $\theta^{(j)}$ (parameter for user j):

$$\min_{\theta^{(j)}} \frac{1}{2} \sum_{i:r(i,j)=1} ((\theta^{(j)})^T (x^{(i)}) - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{k=1}^n (\theta_k^{(j)})^2$$

Optimization algorithm:

$$J(\theta^{(1)}, \dots, \theta^{(n_u)}) = \frac{1}{2} \sum_{j=1}^{n_u} \sum_{i:r(i,j)=1} ((\theta^{(j)})^T (x^{(i)}) - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (x_k^{(j)})^2$$
$$\min_{\theta^{(1)}, \dots, \theta^{(n_u)}} J(\theta^{(1)}, \dots, \theta^{(n_u)})$$

Gradient descent update:

$$\theta_k^{(j)} := \theta_k^{(j)} - \alpha \sum_{i:r(i,j)=1} ((\theta^{(j)})^T (x^{(i)}) - y^{(i,j)}) x_k^{(i)} \quad (\text{for } k = 0)$$

$$\theta_k^{(j)} := \theta_k^{(j)} - \alpha \left(\sum_{i:r(i,j)=1} ((\theta^{(j)})^T (x^{(i)}) - y^{(i,j)}) x_k^{(i)} + \lambda \theta_k^{(j)} \right) \quad (\text{for } k \neq 0)$$

Collaborative filtering

Optimization algorithm:

Given $\theta^{(1)}, \dots, \theta^{(n_u)}$, to learn $x^{(i)}$:

$$\min_{\theta^{(j)}} \frac{1}{2} \sum_{i:r(i,j)=1} ((\theta^{(j)})^T(x^{(i)}) - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{k=1}^n (x_k^{(j)})^2$$

Given $\theta^{(1)}, \dots, \theta^{(n_u)}$, to learn $x^{(1)}, \dots, x^{(n_m)}$:

$$J(\theta^{(1)}, \dots, \theta^{(n_u)}) = \frac{1}{2} \sum_{j=1}^{n_m} \sum_{i:r(i,j)=1} ((\theta^{(j)})^T(x^{(i)}) - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (x_k^{(j)})^2$$

Collaborative filtering algorithm

Given $x^{(1)}, \dots, x^{(n_m)}$ and movie ratings, can estimate $\theta^{(1)}, \dots, \theta^{(n_u)}$:

$$\min_{\theta^{(1)}, \dots, \theta^{(n_u)}} \frac{1}{2} \sum_{j=1}^{n_u} \sum_{i:r(i,j)=1} ((\theta^{(j)})^T(x^{(i)}) - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^{(j)})^2$$

Given $\theta^{(1)}, \dots, \theta^{(n_u)}$, can estimate $x^{(1)}, \dots, x^{(n_m)}$:

$$\min_{x^{(1)}, \dots, x^{(n_m)}} \frac{1}{2} \sum_{j=1}^{n_u} \sum_{i:r(i,j)=1} ((\theta^{(j)})^T(x^{(i)}) - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_m} \sum_{k=1}^n (x_k^{(j)})^2$$

Minimizing $x^{(1)}, \dots, x^{(n_m)}$ and $\theta^{(1)}, \dots, \theta^{(n_u)}$ simultaneously:

$$J(x^{(1)}, \dots, x^{(n_m)}, \theta^{(1)}, \dots, \theta^{(n_u)}) = \frac{1}{2} \sum_{(i,j):r(i,j)=1} ((\theta^{(j)})^T(x^{(i)}) - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_m} \sum_{k=1}^n (x_k^{(j)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^{(j)})^2$$

$$\min_{x, \theta} J(x^{(1)}, \dots, x^{(n_m)}, \theta^{(1)}, \dots, \theta^{(n_u)})$$

Step:

1. Initialize $x^{(1)}, \dots, x^{(n_m)}, \theta^{(1)}, \dots, \theta^{(n_u)}$ to small random values
2. Minimize $J(x^{(1)}, \dots, x^{(n_m)}, \theta^{(1)}, \dots, \theta^{(n_u)})$ using gradient descent
3. For a user with parameters θ and a movie with features x , predict a star rating of $\theta^T x$

Vectorization: Low Rank Matrix Factorization

How to find movies j related to movie i ?

$$\text{small } ||x^{(i)} - x^{(j)}||$$

Implementational Detail: Mean Normalization

recommend a new user the popular(high x) movies.

17.Large Scale Machine Learning

Learning with Large Datasets

Stochastic Gradient Descent

- Batch gradient descent

$$J_{train}(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Repeat{

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

}

- Stochastic gradient descent

$$cost(\theta, (x^{(i)}, y^{(i)})) = \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$J_{train}(\theta) = \frac{1}{m} \sum_{i=1}^m cost(\theta, (x^{(i)}, y^{(i)}))$$

1. Randomly shuffle dataset

2. Repeat{

for i in range(1,m):

$$\theta_j := \theta_j - \alpha (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

}

Mini-Batch Gradient Descent

Batch gradient descent: Use m example in each iteration

Stochastic Gradient Descent: Use 1 example in each iteration

Mini-Batch Gradient Descent: Use b example in each iteration

b = 2~100

Stochastic Gradient Descent Convergence

Online Learning

Map-reduce and data parallelism

Combine:

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^n (temp_j^{(i)})$$

Many learning algorithms can be expressed as computing sums of functions over the training set.

18.Application Example: Photo OCR

Problem Description and Pipeline

1. Text detection
2. Character segmentation
3. Character classification

Sliding Windows

Getting Lots of Data and Artificial Data Synthesis

Ceiling Analysis: What Part of the Pipeline to Work on Next